

L4_GOV_EquitableAccessPricing

Equitable access & pricing policy

For Gemini:

Links:

<https://ai.google.dev/gemini-api/docs/pricing>

<https://cloud.google.com/startup/ai>

https://support.google.com/a/answer/14700766?hl=en&co=DASHER_Family%3DEducation

<https://www.google.com/nonprofits/success-stories/action-against-hunger/>

Gemini Developer API pricing

content_copy

Start building free of charge with generous limits, then scale up with pay-as-you-go pricing for your production ready applications.

Free

For developers and small projects getting started with the Gemini API.

- check_circleLimited access to certain models
- check_circleFree input & output tokens
- check_circleGoogle AI Studio access
- check_circleContent used to improve our products*

[Get started for Free](#)

Paid

For production applications that require higher volumes and advanced features.

- check_circleHigher rate limits for production deployments
- check_circleAccess to Context caching
- check_circleBatch API (50% cost reduction)
- check_circleAccess to Google's most advanced models
- check_circleContent *not* used to improve our products*

[Upgrade to Paid](#)

Enterprise

For large-scale deployments with custom needs for security, support, and compliance, powered by [Vertex AI](#).

- check_circleAll features in Paid, plus optional access to:
- check_circleDedicated support channels
- check_circleAdvanced security & compliance
- check_circleProvisioned throughput
- check_circleVolume-based discounts (based on usage)
- check_circleML ops, model garden and more

[Contact Sales](#)

Gemini 3 Pro Preview

gemini-3-pro-preview

[Try it in Google AI Studio](#)

The best model in the world for multimodal understanding, and our most powerful agentic and vibe-coding model yet.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Not available	\$2.00, prompts <= 200k tokens \$4.00, prompts > 200k tokens
Output price (including thinking tokens)	Not available	\$12.00, prompts <= 200k tokens \$18.00, prompts > 200k
Context caching price	Not available	\$0.20, prompts <= 200k tokens \$0.40, prompts > 200k \$4.50 / 1,000,000 tokens per hour (storage price)

	Free Tier	Paid Tier, per 1M tokens in USD
Grounding with Google Search*	Not available	5,000 prompts per month (free), then (Coming soon) queries
Grounding with Google Maps	Not available	Not available
Used to improve our products	Yes	No

* A customer-submitted request to Gemini may result in one or more queries to Google Search. You will be charged for each individual search query performed.

** Gemini 3 billing for Grounding with Google Search will start January 5, 2026.

Gemini 3 Pro Image Preview 

gemini-3-pro-image-preview

[Try it in Google AI Studio](#)

Our native image generation model, optimized for speed, flexibility, and contextual understanding. **Text input and output** is priced the same as [Gemini 3 Pro](#).

Preview models may change before becoming stable and have more restrictive rate limits.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens
Input price	Not available	\$2.00 (text/image), equivalent to \$0.001

	Free Tier	Paid Tier, per 1M tok
Output price	Not available	\$12.00 (text and thin \$120.00 (images) equivalent to \$0.134 and \$0.24 per 4K ima
Used to improve our products	Yes	No

* Image input is set at 560 tokens or \$0.0011 per image.

** Image output is priced at \$120 per 1,000,000 tokens. Output images from 1024x1024px (1K) and up to 2048x2048px (2K) consume 1120 tokens and are equivalent to \$0.134 per image. Output images up to 4096x4096px (4K) consume 2000 tokens and are equivalent to \$0.24 per image.

Gemini 2.5 Pro

gemini-2.5-pro

[Try it in Google AI Studio](#)

Our state-of-the-art multipurpose model, which excels at coding and complex reasoning tasks.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens in U
Input price	Free of charge	\$1.25, prompts <= 200k token \$2.50, prompts > 200k token
Output price (including thinking tokens)	Free of charge	\$10.00, prompts <= 200k token \$15.00, prompts > 200k
Context caching price	Not available	\$0.125, prompts <= 200k token \$0.25, prompts > 200k \$4.50 / 1,000,000 tokens per

	Free Tier	Paid Tier, per 1M tokens in USD
Grounding with Google Search	Not available	1,500 RPD (free), then \$35 /
Grounding with Google Maps	Not available	10,000 RPD (free), then \$25 /
Used to improve our products	<u>Yes</u>	<u>No</u>

Gemini 2.5 Flash

gemini-2.5-flash

[Try it in Google AI Studio](#)

Our first hybrid reasoning model which supports a 1M token context window and has thinking budgets.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	\$0.30 (text / image / video) \$1.00 (audio)
Output price (including thinking tokens)	Free of charge	\$2.50
Context caching price	Not available	\$0.03 (text / image / video) \$0.1 (audio) \$1.00 / 1,000,000 tokens per hour (storage price)
Grounding with Google Search	Free of charge, up to 500 RPD (limit shared with Flash-Lite RPD)	1,500 RPD (free, limit shared with Flash-Lite RPD)

	Free Tier	Paid Tier, per 1M tokens in USD
Grounding with Google Maps	500 RPD	1,500 RPD (free), then \$25 / 1,000 grounded RPD
Used to improve our products	Yes	No

Gemini 2.5 Flash Preview

gemini-2.5-flash-preview-09-2025

[Try it in Google AI Studio](#)

The latest model based on the 2.5 Flash model. 2.5 Flash Preview is best for large scale processing, low-latency, high volume tasks that require thinking, and agentic use cases.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	\$0.30 (text / image / video) \$1.00 (audio)
Output price (including thinking tokens)	Free of charge	\$2.50
Context caching price	Not available	\$0.03 (text / image / video) \$0.1 (audio) \$1.00 / 1,000,000 tokens per hour (storage price)

	Free Tier	Paid Tier, per 1M tokens in USD
Grounding with Google Search	Free of charge, up to 500 RPD (limit shared with Flash-Lite RPD)	1,500 RPD (free, limit shared with Flash-Lite grounded prompts)
Used to improve our products	Yes	No

Gemini 2.5 Flash-Lite

gemini-2.5-flash-lite

[Try it in Google AI Studio](#)

Our smallest and most cost effective model, built for at scale usage.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens in USD
Input price (text, image, video)	Free of charge	\$0.10 (text / image / video) \$0.30 (audio)
Output price (including thinking tokens)	Free of charge	\$0.40
Context caching price	Not available	\$0.01 (text / image / video) \$0.03 (audio) \$1.00 / 1,000,000 tokens per hour (storage p

	Free Tier	Paid Tier, per 1M tokens in USD
Grounding with Google Search	Free of charge, up to 500 RPD (limit shared with Flash RPD)	1,500 RPD (free, limit shared with Flash RPD) grounded prompts
Grounding with Google Maps	500 RPD	1,500 RPD (free), then \$25 / 1,000 grounded prompts
Used to improve our products	Yes	No

Gemini 2.5 Flash-Lite Preview

gemini-2.5-flash-lite-preview-09-2025

[Try it in Google AI Studio](#)

The latest model based on Gemini 2.5 Flash lite optimized for cost-efficiency, high throughput and high quality.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens in USD
Input price (text, image, video)	Free of charge	\$0.10 (text / image / video) \$0.30 (audio)
Output price (including	Free of charge	\$0.40

	Free Tier	Paid Tier, per 1M tokens in USD
thinking tokens)		
Context caching price	Not available	\$0.01 (text / image / video) \$0.03 (audio) \$1.00 / 1,000,000 tokens per hour (storage price)
Grounding with Google Search	Free of charge, up to 500 RPD (limit shared with Flash RPD)	1,500 RPD (free, limit shared with Flash RPD) grounded prompts
Used to improve our products	<u>Yes</u>	<u>No</u>

Gemini 2.5 Flash Native Audio (Live API)

gemini-2.5-flash-native-audio-preview-09-2025

[Try it in Google AI Studio](#)

Our [Live API](#) native audio models optimized for higher quality audio outputs with better pacing, voice naturalness, verbosity, and mood.

Preview models may change before becoming stable and have more restrictive rate limits.

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	\$0.50 (text) \$3.00 (audio / video)
Output price (including thinking tokens)	Free of charge	\$2.00 (text) \$12.00 (audio)

	Free Tier	Paid Tier, per 1M tok
Used to improve our products	Yes	No

Gemini 2.5 Flash Image 

gemini-2.5-flash-image

[Try it in Google AI Studio](#)

Our native image generation model, optimized for speed, flexibility, and contextual understanding. Text input and output is priced the same as [2.5 Flash](#).

Preview models may change before becoming stable and have more restrictive rate limits.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tok
Input price	Not available	\$0.30 (text / image)
Output price	Not available	\$0.039 per image*
Used to improve our products	Yes	No

[*] Image output is priced at \$30 per 1,000,000 tokens. Output images up to 1024x1024px consume 1290 tokens and are equivalent to \$0.039 per image.

Gemini 2.5 Flash Preview TTS

gemini-2.5-flash-preview-tts

[Try it in Google AI Studio](#)

Our 2.5 Flash text-to-speech audio model optimized for price-performant, low-latency, controllable speech generation.

Preview models may change before becoming stable and have more restrictive rate limits.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tok
Input price	Free of charge	\$0.50 (text)
Output price	Free of charge	\$10.00 (audio)
Used to improve our products	<u>Yes</u>	<u>No</u>

Gemini 2.5 Pro Preview TTS

gemini-2.5-pro-preview-tts

[Try it in Google AI Studio](#)

Our 2.5 Pro text-to-speech audio model optimized for powerful, low-latency speech generation for more natural outputs and easier to steer prompts.

Preview models may change before becoming stable and have more restrictive rate limits.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tok
Input price	Not available	\$1.00 (text)
Output price	Not available	\$20.00 (audio)
Used to improve our products	<u>Yes</u>	<u>No</u>

Gemini 2.0 Flash

gemini-2.0-flash

[Try it in Google AI Studio](#)

Our most balanced multimodal model with great performance across all tasks, with a 1 million token context window, and built for the era of Agents.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens in LLM
Input price	Free of charge	\$0.10 (text / image / video) \$0.70 (audio)
Output price	Free of charge	\$0.40
Context caching price	Free of charge	\$0.025 / 1,000,000 tokens (\$0.175 / 1,000,000 tokens (
Context caching (storage)	Not available	\$1.00 / 1,000,000 tokens per
Image generation pricing	Free of charge	\$0.039 per image*
Tuning price	Not available	Not available
Grounding with Google Search	Free of charge, up to 500 RPD	1,500 RPD (free), then \$35 ,
Grounding with Google Maps	500 RPD	1,500 RPD (free), then \$25 ,
Used to improve our products	<u>Yes</u>	<u>No</u>

[*] Image output is priced at \$30 per 1,000,000 tokens. Output images up to 1024x1024px consume 1290 tokens and are equivalent to \$0.039 per image.

Gemini 2.0 Flash-Lite

gemini-2.0-flash-lite

[Try it in Google AI Studio](#)

Our smallest and most cost effective model, built for at scale usage.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens in LLM
Input price	Free of charge	\$0.075

	Free Tier	Paid Tier, per 1M tok
Output price	Free of charge	\$0.30
Context caching price	Not available	Not available
Context caching (storage)	Not available	Not available
Tuning price	Not available	Not available
Grounding with Google Search	Not available	Not available
Used to improve our products	Yes	No

Imagen 4

imagen-4.0-generate-001, imagen-4.0-ultra-generate-001, imagen-4.0-fast-generate-001

[Try it in Google AI Studio](#)

Our latest image generation model, with significantly better text rendering and better overall image quality.

Preview models may change before becoming stable and have more restrictive rate limits.

	Free Tier	Paid Tier, per Image
Imagen 4 Fast image price	Not available	\$0.02
Imagen 4 Standard image price	Not available	\$0.04
Imagen 4 Ultra image price	Not available	\$0.06
Used to improve our products	Yes	No

Imagen 3

imagen-3.0-generate-002

[Try it in Google AI Studio](#)

Our state-of-the-art image generation model, available to developers on the paid tier of the Gemini API.

	Free Tier	Paid Tier, per Image
Image price	Not available	\$0.03
Used to improve our products	Yes	No

Veo 3.1

veo-3.1-generate-preview, veo-3.1-fast-generate-preview

[Try Veo 3.1](#)

Our latest video generation model, available to developers on the paid tier of the Gemini API.

Preview models may change before becoming stable and have more restrictive rate limits.

	Free Tier	Paid Tier, per second
Veo 3.1 Standard video with audio price (default)	Not available	\$0.40
Veo 3.1 Fast video with audio price (default)	Not available	\$0.15
Used to improve our products	Yes	No

Note: In some cases, an audio processing issue may prevent a video from being generated.

You will only be charged if your video is successfully generated.

Veo 3

veo-3.0-generate-001, veo-3.0-fast-generate-001

[Try Veo 3](#)

Our stable video generation model, available to developers on the paid tier of the Gemini API.

	Free Tier	Paid Tier, per second
Veo 3 Standard video with audio price (default)	Not available	\$0.40
Veo 3 Fast video with audio price (default)	Not available	\$0.15
Used to improve our products	Yes	No

Note: In some cases, an audio processing issue may prevent a video from being generated. You will only be charged if your video is successfully generated.

Veo 2

veo-2.0-generate-001

[Try the API](#)

Our state-of-the-art video generation model, available to developers on the paid tier of the Gemini API.

	Free Tier	Paid Tier, per second
Video price	Not available	\$0.35
Used to improve our products	Yes	No

Gemini Embedding

gemini-embedding-001

[Try the API](#)

Our newest embeddings model, more stable and with higher rate limits than previous versions, available to developers on the free and paid tiers of the Gemini API.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens
Input price	Free of charge	\$0.15
Used to improve our products	Yes	No

Gemini Robotics-ER 1.5 Preview

gemini-robotics-er-1.5-preview

[Try it in Google AI Studio](#)

Gemini Robotics-ER, short for Gemini Robotics-Embodied Reasoning, is a thinking model that enhances robots' abilities to understand and interact with the physical world.

[StandardBatch](#)

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	\$0.30 (text / image / video) \$1.00 (audio)
Output price (including thinking tokens)	Free of charge	\$2.50
Grounding with Google Search	Free of charge, up to 500 RPD (limit shared with Flash-Lite RPD)	1,500 RPD (free, limit shared with Flash-Lite grounded prompts)
Used to improve our products	Yes	No

Gemini 2.5 Computer Use Preview

gemini-2.5-computer-use-preview-10-2025

Our Computer Use model optimized for building browser control agents that automate tasks.

	Free Tier	Paid Tier, per 1M tok
Input price	Not available	\$1.25, prompts <= 200 \$2.50, prompts > 200
Output price	Not available	\$10.00, prompts <= 200 \$15.00, prompts > 200
Used to improve our products	Yes	No

Gemma 3

Our lightweight, state-of the art, open model built from the same technology that powers our Gemini models.

	Free Tier	Paid Tier, per 1M tok
Input price	Free of charge	Not available
Output price	Free of charge	Not available
Context caching price	Free of charge	Not available
Context caching (storage)	Free of charge	Not available
Tuning price	Not available	Not available
Grounding with Google Search	Not available	Not available
Used to improve our products	Yes	No

Gemma 3n

Our open model built for efficient performance on everyday devices like mobile phones, laptops, and tablets.

	Free Tier	Paid Tier, per 1M tokens
Input price	Free of charge	Not available
Output price	Free of charge	Not available
Context caching price	Free of charge	Not available
Context caching (storage)	Free of charge	Not available
Tuning price	Not available	Not available
Grounding with Google Search	Not available	Not available
Used to improve our products	Yes	No

Pricing for Tools

Tools are priced at their own rates, applied to the model using them. Check the [Models](#) page for which tools are available to each model.

	Free Tier	Paid Tier, per 1M tokens in USD
Google Search	500 RPD free (limit shared for Flash and Flash-Lite). Not available for Pro.	1,500 RPD free (limit shared for Flash and Flash-Lite). Then \$35 / 1,000 grounded prompts See the Gemini 3 Pro Preview pricing section for more pricing.
Google Maps	500 RPD Not available for Pro.	1,500 RPD free (limit shared for Flash and Flash-Lite). 10,000 RPD free for Pro. Then \$25 / 1,000 grounded prompts

	Free Tier	Paid Tier, per 1M tokens in USD
<u>Code execution</u>	Free of charge	Free of charge
<u>URL context</u>	Free of charge	Charged as input tokens per model pricing.
<u>Computer use</u>	Not available	See Gemini 2.5 Computer Use Preview pricing table
<u>File search</u>	Free of charge	Charged for embeddings at \$0.15 / 1M tokens. Retrieved document tokens charged as regular tokens

[*] Google AI Studio usage is free of charge in all [available regions](#). See [Billing FAQs](#) for details.

[**] Prices may differ from the prices listed here and the prices offered on Vertex AI. For Vertex prices, see the [Vertex AI pricing page](#).

[***] If you are using [dynamic retrieval](#) to optimize costs, only requests that contain at least one grounding support URL from the web in their response are charged for Grounding with Google Search. Costs for Gemini always apply. Rate limits are subject to change.

Our AI program is built to accelerate AI Startups

Receive \$350,000 in Cloud credits, resources to accelerate your startup journey, dedicated support and enablement, and the best of Google's open AI ecosystem.

A program designed for Scale tier AI startups

Take the next step in your AI journey

Get the technology, community, and resources you need to build and grow your AI startup faster and easier.

Additional resources to build with Google AI

Our most powerful agentic and vibe-coding model yet is now available to every developer and enterprise team. Transform your businesses for the agentic future.

- State-of-the-art reasoning and multimodality
- Powerful agentic coding and front-end creation
- Advanced tool use and planning

Get the knowledge needed to go from an idea to prototype to scale, whether your goals are to automate tasks, enhance creativity, or launch entirely new user experiences for your startup.

- Maximize efficiency and productivity
- Accelerate your product development
- Leverage multimodal capabilities with Gemini

Join us at Google Cloud Next '26 on April 22-24, 2026 at Mandalay Bay Convention Center, Las Vegas to stay ahead of the curve and connect with AI innovators.

- Dive into your favorite topics
- Meet with AI experts
- Create new possibilities
- Join us in the Startup Hub for startup-specific sessions and networking opportunities

Compare Gemini for Google Workspace add-ons

Google Workspace with Gemini add-ons enable organizations to use Google's generative AI capabilities with many Google Workspace editions. Learn availability, and compare add-ons in these categories:

- **Gemini for Business / Enterprise—Legacy**—Note: these add-ons are longer available for purchase. [Learn more](#).
- **Gemini for Education**—AI capabilities for qualifying educational institutions

Education Business / Enterprise

Compare Gemini for Education add-ons

Use these tables to compare generative AI features across the following add-ons:

Google AI Pro for Education—Generative AI with full access and usage, and additional features for AI-powered meetings and security protections

Learn more about [Gemini for Google Workspace](#).

On this page

- [Add-on availability](#)
- [Gemini pricing](#)
- [Gemini access](#)
- [Gemini in Meet](#)
- [Gemini in Gmail](#)
- [Gemini in Docs editors](#)
- [Gemini in Chat](#)
- [Additional Gemini features](#)

Add-on availability

Use this table to find out which Google Workspace editions support the add-ons shown below.

Compare instead: [Business / Enterprise add-ons](#)

Google Workspace edition	Google AI Pro for Education*
Education Fundamentals	✓
Education Standard	✓
Education Plus	✓

* Subject to [monthly usage limits](#)

[↑ back to top](#)

Gemini pricing

Compare instead: [Business / Enterprise add-ons](#)

The following chart shows pricing for Gemini for Google Workspace. Standard prices are shown. Contact your reseller for information about available discounts.

	Flexible Plan	Annual/Fixed-Term Plan
Monthly payment	Google AI Pro for Education: \$24 USD per user	Google AI Pro for Education: \$20 USD per user

[↑ back to top](#)

Gemini app access

The Gemini app is available as a Core service for Google Workspace for Education users of all ages with [enterprise-grade data protection](#), meaning your chats will not be reviewed by anyone, used to train generative machine-learning technologies, or shared with other users or organizations. Your chats will still be subject to the [Google Workspace for Education terms of service](#).

Gemini in Meet

Compare instead: [Business / Enterprise add-ons](#)

	Google AI Pro for Education
Create custom backgrounds	✓
Studio look to improve low quality video	✓
Studio lighting to simulate professional lighting	✓
Studio sound to improve audio	✓
Live translated captions in 65+ languages	✓
Take notes for me *	✓
Adaptive audio	✓

*Requires turning on [Gemini Alpha](#) features

[↑ back to top](#)

Gemini in Gmail

Compare instead: [Business / Enterprise add-ons](#)

	Google AI Pro for Education
Help me write emails *	✓
Side panel	✓
Salesforce in Gemini extension	✓

* Available in [English, Spanish, and Portuguese](#)

[↑ back to top](#)

Gemini in Docs editors

Compare instead: [Business / Enterprise add-ons](#)

Google AI Pro for Education

Help me write (Docs) **	✓
Proofread (Docs)	✓
Enhanced Smart Fill (Sheets)	✓
Help me visualize (image creation in Slides)	✓
Remove image background (Slides)	✓
Side Panel (Docs , Sheets , Slides , and Drive)	✓
Help me visualize (image creation in Google Docs)	✓

* Available in [English](#), [Spanish](#), and [Portuguese](#)

[↑ back to top](#)

Gemini in Chat

Compare instead: [Business / Enterprise add-ons](#)

Google AI Pro for Education

Gemini Summaries in Google Chat	✓
Side panel	✓

[↑ back to top](#)

Additional Gemini features

Compare instead: [Business / Enterprise add-ons](#)

Google AI Pro for Education

Google Vids	✓
NotebookLM	✓

* Requires an admin to turn on [Gemini Alpha](#) features